

★ 1. What is Logistic Regression?

Logistic Regression is a **classification algorithm**, NOT a regression model (despite the name).

- ✓ Predicts **categories** → Example: Spam / Not Spam
 - ✓ Outputs **probabilities** between **0 and 1**
 - ✓ Converts probability into class label (0 or 1) using a **threshold**
-

★ 2. Why do we call it Regression?

Because it uses a linear model:

$$z = w_1x_1 + w_2x_2 + \dots + b$$

But applies the **sigmoid function** on top to convert $z \rightarrow$ probability.

★ 3. Sigmoid Function

The heart of Logistic Regression.

$$p = \frac{1}{1 + e^{-z}}$$

- ✓ Output always lies between **0 and 1**
- ✓ Used to interpret model output as “probability of class 1”

Example:

$$\begin{aligned} z = 2.0 &\rightarrow \text{sigmoid} = 0.88 \rightarrow 88\% \text{ chance of class 1} \\ z = -1.0 &\rightarrow \text{sigmoid} = 0.27 \rightarrow 27\% \text{ chance of class 1} \end{aligned}$$

★ 4. Probability → Class Conversion (Threshold)

The model does NOT directly predict 0 or 1.

It predicts probability p .

We use a threshold (default = 0.5) to convert:

```
if p ≥ 0.5 → class = 1  
if p < 0.5 → class = 0
```

Example:

```
p=0.84 → 1  
p=0.33 → 0  
p=0.51 → 1
```

★ 5. Why is Threshold important?

You can modify threshold based on domain:

- ✓ Medical diagnosis → need high recall → threshold low (0.3)
- ✓ Fraud detection → avoid false alarms → threshold high (0.7)
- ✓ Normal binary classification → threshold = 0.5

ROC curve helps identify the **best threshold**.

★ 6. Loss Function Used: Log Loss (Cross-Entropy)

Logistic Regression tries to minimize:

$$\text{Loss} = -[y \log(p) + (1-y) \log(1-p)]$$

- ✓ Punishes confident wrong predictions
 - ✓ Improves probability calibration
-

★ 7. Regularization in Logistic Regression

To prevent overfitting:

✓ L2 (Ridge Regularization)

- Shrinks weights
- Does NOT make coefficients zero
- DEFAULT penalty
- Very stable
- Works with most solvers

✓ L1 (Lasso Regularization)

- Makes some weights exactly **zero**
- Performs **feature selection**
- Works with solvers: `liblinear`, `saga`

✓ ElasticNet (L1 + L2 combination)

- Balances sparsity and stability
 - Works only with solver: `saga`
-

★ 8. Solvers (Optimization Algorithms)

Solver	Supports	Use Case
lbfgs	L2	Best default, fast
newton-c g	L2	Multiclass problems
liblinear	L1, L2	Small datasets
sag	L2	Very large datasets
saga	L1, L2, Elastic	Best for large sparse datasets

★ 9. Confusion Matrix

Used to compute TPR, FPR, Accuracy, etc.

	Pred 1	Pred 0
Actual 1	TP	FN
Actual 0	FP	TN

Definitions:

- **TP:** Correct positive
 - **TN:** Correct negative
 - **FP:** Wrong positive (false alarm)
 - **FN:** Missed positive
-

★ 10. Important Metrics

=====

LOGISTIC REGRESSION METRICS

=====

1. Accuracy

$$\text{Accuracy} = (\text{TP} + \text{TN}) / (\text{TP} + \text{TN} + \text{FP} + \text{FN})$$

Accuracy tells how many total predictions were correct.
Fails on imbalanced datasets.

2. Precision

$$\text{Precision} = \text{TP} / (\text{TP} + \text{FP})$$

Meaning:

When the model predicts Positive, how often is it correct?
Avoids False Positives.

3. Recall (TPR) - True Positive Rate

$$\text{Recall} = \text{TP} / (\text{TP} + \text{FN})$$

Meaning:

How many actual positives did the model catch?
Avoids False Negatives.

4. FPR - False Positive Rate

$$\text{FPR} = \text{FP} / (\text{FP} + \text{TN})$$

Meaning:

Out of all actual negatives, how many were wrongly predicted as positive?

5. Confusion Matrix

Predicted

| 1 | 0

	1	0
Actual 1	TP	FN
Actual 0	FP	TN

Where:

TP = True Positive

TN = True Negative

FP = False Positive

FN = False Negative

6. F1 Score

$$F1 = 2 * (\text{Precision} * \text{Recall}) / (\text{Precision} + \text{Recall})$$

Balances Precision and Recall.

Useful for imbalanced datasets.

7. ROC Curve

ROC curve plots:

X-axis → FPR

Y-axis → TPR (Recall)

Shows performance at different probability thresholds.

8. AUC Score

AUC = Area Under ROC Curve

AUC = 1 → perfect model

AUC = 0.5 → random guessing

AUC < 0.5 → worse than random

=====

END OF LOGISTIC METRICS NOTES

=====

★ 11. ROC Curve (Receiver Operating Characteristic)

★ 11. ROC Curve (Receiver Operating Characteristic)

ROC curve plots:

- **X-axis:** FPR
- **Y-axis:** TPR

Purpose:

- ✓ Shows model performance for **all thresholds**
- ✓ Helps visualize separability power

HOW ROC IS MADE?

For every threshold:

1. Convert probability → predicted class
2. Compute TP, FP, FN, TN
3. Compute TPR & FPR
4. Plot point (FPR, TPR)

ROC curve = all such points.

★ 12. AUC Score (Area Under Curve)

Measures total area under ROC curve.

AUC Score	Model Quality
-----------	---------------

1.0	Perfect
0.9–1.0	Excellent
0.8–0.9	Good

0.7–0.8	Okay
0.5	Random guessing
< 0.5	Worse than random

AUC = probability model ranks positive higher than negative.

⭐ 13. Understanding Thresholds on ROC Graph

From code:

```
for xyz in zip(model_fpr, model_tpr, thresholds):  
    ax.annotate(str(round(xyz[2],2)), xy=(xyz[0], xyz[1]))
```

Meaning:

- `xyz[0]` = FPR
- `xyz[1]` = TPR
- `xyz[2]` = threshold

This writes the **threshold value** at each ROC point.

⭐ 14. Why ROC-AUC is Better Than Accuracy

Accuracy fails when data is **imbalanced**.

Example:

- 99% class 0
- 1% class 1

Model predicts all zeros → 99% accuracy but **useless**.

ROC-AUC measures **ranking quality**, not class labels.

★ 15. Key Advantages of Logistic Regression

- ✓ Easy to understand
 - ✓ Fast
 - ✓ Works on large datasets
 - ✓ Outputs probability (very useful)
 - ✓ Regularization available
 - ✓ Performs well if classes are linearly separable
-

★ 16. Limitations

- ✗ Cannot capture complex non-linear relationships
 - ✗ Performance drops if features not scaled
 - ✗ Poor on heavily imbalanced data (without AUC/threshold tuning)
-

★ 17. Example Summary (Simple)

Suppose probabilities:

[0.95, 0.80, 0.40, 0.30]

Threshold = 0.8

Predictions:

[1, 1, 0, 0]

Confusion matrix → compute TPR, FPR → plot ROC point.

⭐ 18. When to use Logistic Regression?

Use when:

- ✓ Binary classification
 - ✓ Features are linearly separable
 - ✓ You need interpretability
 - ✓ You want probability output
 - ✓ Dataset is not too huge or too complex
-

⭐ 19. Important Notes About Thresholds

- Threshold = 0.5 is **NOT always best**
 - Use ROC-AUC to choose best threshold
 - Use precision–recall curve for imbalance
 - Domain-specific thresholds improve results
-

⭐ 20. Best Practice Pipeline

1. Train Logistic Regression with scaling
2. Tune hyperparameters (solver, C, penalty)
3. Plot ROC & compute AUC
4. Select best threshold
5. Evaluate with confusion matrix
6. Deploy model with chosen threshold



ROC Curve – Notes

1. Logistic Regression gives probabilities

Example:

[0.95, 0.80, 0.40, 0.30]

These are NOT classes — they are probabilities of being class 1.

2. Converting Probability → Class (0/1)

To make predictions:

```
if probability ≥ threshold → predict 1  
else → predict 0
```

Default threshold = 0.5

But ROC does **not** use only one threshold — it tries many.

3. How ROC Curve Works (Step-by-Step)

ROC tests **all possible thresholds**, for example:

[1.0, 0.95, 0.80, 0.40, 0.30, 0.0]

For each threshold:

1. Convert probabilities → 0/1 predictions
2. Compute:
 - **TPR (Recall)** = TP / (TP + FN)
 - **FPR** = FP / (FP + TN)
3. Plot the point → (FPR, TPR)

All points combined create the **ROC Curve**.

4. Why We Use ROC Curve?

- Works across **all thresholds**, not just 0.5
 - Best for **imbalanced datasets**
 - Shows model's ability to **separate** classes
 - Helps compare classifiers clearly
-

5. AUC Score (Area Under Curve)

- AUC = 1 → Perfect model
- AUC = 0.5 → Random guessing
- AUC < 0.5 → Worse than random

Bigger AUC = Better model.

6. ROC Graph Axes

- **X-axis** = FPR (False Positive Rate)
 - **Y-axis** = TPR (True Positive Rate / Recall)
-

7. Key Idea (One Line)

ROC changes the threshold again and again, calculates TPR & FPR, plots all points → this becomes the ROC curve.



Hyperparameter Tuning Notes

GridSearchCV vs RandomizedSearchCV

★ 1. Why Hyperparameter Tuning?

Machine learning models have **hyperparameters** (like `C`, `penalty`, `solver`, `max_depth`, `n_estimators`)
that control learning.

Bad hyperparameters → poor accuracy

Good hyperparameters → high accuracy

Hyperparameter tuning finds the best combination automatically.

◆ GRIDSEARCHCV

✓ Definition

GridSearchCV tests every possible combination of hyperparameters.

It performs:

- Exhaustive search
 - Cross-validation for each combination
 - Picks the combination with the highest score
-

✓ How GridSearchCV Works (Internally)

Given:

```
C = [0.1, 1, 10]
penalty = ['l1', 'l2']
```

Total combinations = $3 \times 2 = 6$

For each combination:

1. Perform K-fold cross-validation
2. Compute mean accuracy
3. Store the score
4. Move to next combination

Finally selects `best_params_`, `best_score_`, `best_estimator_`

✓ Advantages

- Guaranteed to find the best hyperparameter set
- Simple & exhaustive
- Works well when parameter search space is small

✓ Disadvantages

- Slow
 - Expensive for large parameter grids
 - Not good for high-dimensional spaces
-



RANDOMIZEDSEARCHCV

✓ Definition

RandomizedSearchCV tests only a random subset of the hyperparameter combinations.

Instead of trying all combinations, it tries `n_iter` random ones.

✓ How RandomizedSearchCV Works (Internally)

You tell it:

```
n_iter = 20
```

It will:

1. Randomly sample 20 combinations

2. For each combination → run cross-validation
3. Pick the best one among the sampled options

It does not try all combinations.

✓ Advantages

- Much faster
- Works well with large hyperparameter spaces
- Can sample from distributions (uniform, log-uniform, etc.)
- Good when model training is expensive (e.g., XGBoost, Random Forest, Neural Nets)

✓ Disadvantages

- Does not guarantee the absolute best combination
 - Performance depends on number of iterations
-

◆ GRIDSEARCHCV vs RANDOMIZEDSEARCHCV (Summary Table)

Feature	GridSearchCV	RandomizedSearchCV
Tries all combinations	✓ Yes	✗ No
Speed	✗ Slow	✓ Fast

Best params guaranteed	✓ Yes	✗ Not guaranteed
Good for small search space	✓	✗
Good for large search space	✗	✓
Supports distributions	✗	✓
Typical use case	Linear/Logistic Regression, SVM small grids	Random Forest, XGBoost, Deep Learning

❖ When To Use Which?

✓ Use GridSearchCV when:

- You have few hyperparameters
- Search space is small
- You need exact optimal values

✓ Use RandomizedSearchCV when:

- Hyperparameter space is huge
 - Model training is slow
 - You want results fast
 - You're exploring parameter ranges
-



Best Practice (Use Both Together)

1. RandomizedSearchCV first
→ Quickly finds good range of hyperparameters
2. GridSearchCV next
→ Fine-tunes within that good range

Best of both worlds:

- ✓ Speed
- ✓ Accuracy