



Decision Trees – Deep Understanding

A **decision tree** is a model that makes decisions by asking questions step-by-step.

Real-life example:

You want to decide **whether a student will pass or fail**.

You may ask questions like:

1. Did they study more than 3 hours daily?
2. Is attendance $\geq 75\%$?
3. Did they submit assignments?

Each question splits students into groups.



Goal:

Create groups where students are **clearly separated** (pure groups).

- One group mostly **Pass**
 - One group mostly **Fail**
-



THE KEY IDEA

Decision Tree wants **PURITY**.

Pure Group

All are same class

Impure Group

Mixed classes

Easy to decide

Confusing

Perfect model prediction

Many errors

Example:

Group A: [Pass, Pass, Pass, Pass] \rightarrow Pure (good)

Group B: [Pass, Fail, Pass, Fail] → Impure (bad)

So the whole game is:

? Which question gives pure groups fastest?

To measure purity/impurity, we use:

1. **Entropy**
 2. **Gini Impurity**
 3. **Information Gain** — decides the best question
-



What is Entropy (in depth)

REAL MEANING OF ENTROPY:

Entropy measures **CONFUSION / DISORDER / UNCERTAINTY**.

Example:

Class of 10 students choosing a game:

Case A:

```
10 want cricket  
0 want football
```

Everyone agrees → No confusion → **Entropy = 0**

Case B:

```
5 want cricket  
5 want football
```

Total conflict, high confusion → **Entropy = High**

Case C:

9 want cricket

1 wants football

Very little confusion → **Entropy = Low**

💡 So entropy is highest when 50-50

Because confusion is maximum.

💡 Entropy is zero when all same

🎯 What is Gini Impurity (in depth)

👉 Gini tells us the probability of making a wrong prediction if we randomly label an item according to the class distribution.

- If the group is **pure** (all same class) → mistake = 0 → **Gini = 0**
- If the group is **mixed** → mistakes high → **Gini high**

🎯 Gini Formula

$$\text{Gini} = 1 - \sum (p_i^2)$$

Where:

- p_{i,p_i} = probability of class i
Only two steps:
 1. Calculate probability of each class
 2. Square them, add them, subtract from 1

🧪 Example 1: Mixed group

Suppose target column Y:

[1, 1, 1, 0, 0, 1, 0, 1] (8 samples)

Step 1: Count classes

Count of 1 = 5.. Count of 0 = 3

Total = 8

Step 2: Probabilities

$p(1) = 5/8 = 0.625$

$p(0) = 3/8 = 0.375$

Step 3: Apply Gini formula $\text{Gini} = 1 - (0.625^2 + 0.375^2)$

$$= 1 - (0.3906 + 0.1406) = 1 - 0.5312 = 0.4688 = 1 - (0.3906 + 0.1406) = 1 - 0.5312 = 0.4688 = 1 - (0.3906 + 0.1406) = 1 - 0.5312 = 0.4688$$

 Interpretation

Gini $\approx 0.47 \rightarrow$ medium impurity (mixed group)

Example 2: Pure group

[1, 1, 1, 1] (4 samples)

Probabilities

$p(1) = 4/4 = 1.0$

$p(0) = 0/4 = 0.0$

Apply formula

$$\text{Gini} = 1 - (1^2 + 0^2)$$

$$= 1 - 1$$

$$= 0$$

✨ Pure group → Gini = 0 → Best possible



Example 3: 50-50 perfect confusion

[1, 1, 0, 0] (4 samples)

Probabilities

$$p(1) = 2/4 = 0.5$$

$$p(0) = 2/4 = 0.5$$

Apply formula

$$\text{Gini} = 1 - (0.5^2 + 0.5^2)$$

$$= 1 - (0.25 + 0.25)$$

$$= 1 - 0.50$$

$$= 0.50$$

💥 Maximum confusion case for two classes.



SUMMARY TABLE

Data Example	Probabilities	Gini	Purity
--------------	---------------	------	--------

[1,1,1,1]	$p_1 = 1, p_0 = 0$	0	Pure (Good)
-----------	--------------------	---	----------------

[1,1,0,0]	$p_1 = .5, p_0 = .5$	0.5	Very Impure
-----------	----------------------	-----	----------------

[1,1,1,0,0,1,0 ,1]	$p_1=.625,$ $p_0=.375$	0.4 7	Medium
-----------------------	---------------------------	----------	--------

REAL MEANING OF GINI:

Gini measures:

Chance that a randomly picked example will be incorrectly classified

Imagine a box:

- 4 apples
- 4 oranges

If you randomly guess "apple"

→ 50% chance of mistake

So **mistake chance is high** → Gini high

If the box has:

- 8 apples, 0 oranges

Random guess "apple"

→ 0% chance of mistake

So **Gini = 0 (perfect purity)**

KEY INTUITION:

- | Very mixed group → High Gini → Bad split |
 - | Pure group → Low Gini → Good split |
-

Entropy vs Gini (deep intuition)

Feature	Entropy	Gini
Meaning	confusion	mistake rate
Curve	complex (log function)	smooth, simple
Calculation	slower	faster
Used in	ID3, C4.5	CART, Random Forest
Best for	theory	real-life speed

Both measure purity

They usually produce similar trees.

What Information Gain Does

 Information Gain tells the decision tree which feature (question) is the BEST to split the data at a node.

It compares different features and says:

"Which feature reduces confusion the most?"

So,

Information Gain = How much confusion is removed by splitting using that feature



REAL-LIFE EXAMPLE (Super Easy)

Imagine you want to separate students into Pass and Fail groups.

You have two possible questions:

Feature S1 → ("Did they attend class?")

Feature S2 → ("Did they study more than 3 hours?")

If you split using S1, students still look mixed:

Group A: [Pass, Fail, Pass]

Group B: [Fail, Pass]

Still confusing! 🤔

If you split using S2, students become very clear:

Group A: [Pass, Pass, Pass, Pass]

Group B: [Fail, Fail]

Perfect clarity! 😍

So S2 is a better question.

Information Gain calculates this difference mathematically.



What is happening in your screenshot

You calculated:

$\text{Gain}(S1) = 0.044$

$\text{Gain}(S2) = 0.094$

That means:

- S2 removes more confusion than S1
- S2 creates purer groups
- S2 is better for splitting data

So the tree will use S2 as the root question.

WHY DO WE EVEN NEED INFORMATION GAIN?

Because without it:

- The tree won't know which feature to ask first
- It may choose a wrong feature
- It may continue making bad splits
- Model accuracy becomes poor

Using Information Gain:

- ✓ Decision Tree asks best question FIRST
 - ✓ Tree becomes short & accurate
 - ✓ Model learns meaningful rules
 - ✓ Better classification and generalization
-

🌟 KEY IDEA (one line)

Information Gain tells us:

How much cleaner the data becomes after splitting on a feature

🌟 Visual Intuition

Before split:

[Pass, Fail, Pass, Fail, Fail, Pass] → Confused, mixed

After split using S2:

Left: [Pass, Pass, Pass] (pure)

Right: [Fail, Fail, Fail] (pure)

Confusion removed → IG high → BEST split

🎁 MEMORY HACK

Entropy = Confusion

Information Gain = Confusion removed

Goal = Choose feature that removes maximum confusion

🧪 In your diagram

The formula:

$$\text{Gain}(S, f) = H(C) - \sum (|S_v| / |S|) * H(S_v)$$

$H(C)$ = Entropy of the parent node

S_v = Subset of S after splitting by feature f

$|S_v|$ = Number of samples in subset S_v

$|S|$ = Number of samples in parent set S

$H(S_v)$ = Entropy of subset S_v

Σ = Sum over all splits (children)

Just means:

Information Gain = Confusion before split – confusion after split



FINAL ANSWER (most important)

Information Gain decides which feature to use to split the data by measuring which one reduces the impurity the most.



Your Example Calculation (Explained)

For two features S_1 and S_2 :

Feature	Info Gain
---------	-----------

S_1	0.044
-------	-------

S_2	0.094
-------	-------

Meaning:

- S_2 removes more confusion
- S_2 creates purer groups
- So S_2 is better and selected for splitting

 **Higher IG = Best Feature**



WHY WE NEED ALL OF THIS?

Because:

- ML must learn rules that separate classes clearly
- Without purity → decision trees overfit & perform poorly
- Selecting best feature improves accuracy and reduces depth

Used in:

Area	Why Important
Random Forest	many trees splitting → needs fast purity check
XGBoost	uses gain to choose split
Explainable AI	tells which feature makes biggest impact
Medical diagnosis	separating healthy vs diseased
Credit approval	separating risky vs safe customers
Fraud detection	separating fraud vs genuine



INTERVIEW QUESTIONS (must know)

Question	Short Answer
What is entropy?	Confusion measure
What is Gini?	Chance of misclassification
What is IG?	Reduction in confusion

Why choose Gini?	Faster computation
Why choose entropy?	better theory separation
Which feature is chosen?	highest IG



Decision Trees Overfitting Problem

Decision trees try to make groups as pure as possible.
To do this, the tree keeps splitting until:

- Every leaf contains a single sample
- Or all samples in a leaf belong to same class

This creates a **very deep and complex tree**.

Problem:

- Tree **fits training data perfectly** (100% accuracy)
- But **fails on new/unseen data** (low test accuracy)

This problem is called **Overfitting**.



Solution = Pruning

Pruning means **cutting off unnecessary branches** of the tree.

Goal:

Make the tree **simpler, generalize better**, and **reduce overfitting**.

✖ Two Types of Pruning

Type	When done?	Meaning
Pre-Pruning (Early Stopping)	Before tree grows fully	Stop splitting early
Post-Pruning	After full tree built	Build full tree, then cut bad branches

1 Pre-Pruning (Prevent overfitting early)

Here we **stop splitting a node early** if the split does not significantly improve purity.

Conditions used to stop splitting:

Rule	Meaning
Min samples split	Stop if fewer than N samples
Min leaf size	Do not create very small leaves
Max depth	Stop after reaching depth limit
Max nodes	Do not create too many nodes
Min information gain	Only split if IG > threshold

Example:

If a node has only 3 samples left:

[1, 1, 0]

Splitting further doesn't help → Stop

Advantages:

- Faster training

- Smaller tree
- Less memory

Disadvantages:

- May stop too early
 - Might miss useful patterns
-

2 Post-Pruning (Cut branches later)

Here we:

1. **Grow full tree completely** (allow overfitting)
2. **Remove branches that do not improve performance on validation data**

How post-pruning works:

Step	Explanation
Build full tree	Even if leaves become extremely deep
Check validation accuracy	Or use cost complexity
Remove weakest branches	Reduce complexity
Keep pruning until accuracy stops improving	Final tree is optimal

Example:

A leaf:

[1, 1, 1, 0]

Splitting into:

Left: [1,1] Right:[1,0]

Makes no real improvement → prune (remove split)

Advantages:

- Best performance
- More accurate tree
- Removes harmful noisy splits

Disadvantages:

- Costly / slower
-



Post-Pruning Algorithms

Algorithm	How it works
Reduced Error Pruning	Remove node if validation accuracy improves
Cost Complexity Pruning (CCP)	Used in CART, controls complexity with α
Minimum Description Length	Uses compression theory

In sklearn:

```
DecisionTreeClassifier(ccp_alpha=0.01)
```



Difference Summary

Feature	Pre-Pruning	Post-Pruning
Time	Fast	Slow

Tree growing	Stops early	Grows fully
Complexity	Less	High after growing
Risk	Underfitting	Better balance
Preferred	Fast needs	Better accuracy

WHAT DOES VARIANCE REDUCTION MEAN?

👉 Variance Reduction means how much improvement we get when we split the data.

More specifically:

How much the spread (difference) between numbers reduces after splitting.

If values are far apart → high variance → unpredictable

If values become close → low variance → easy to predict

So:

Variance Reduction = How much chaos we removed

REAL-LIFE EXAMPLE (simple)

Imagine marks of students:

[30, 95, 40, 98]

These numbers are **very different** → hard to guess next mark.

Variance is **high** (big spread) → like noisy data.

Now suppose we split based on study hours:

After Split:

Group A: [30, 40] → very close → low variance
Group B: [95, 98] → very close → low variance

This split helped us because:

- Each group has similar values
- Prediction becomes easier

So we say:

Variance reduced a lot = Good split



KEY IDEA

Variance Reduction shows how much better (cleaner) our groups became after splitting.

Before Split	After Split
Mixed values	Values close to each other
High variance	Low variance
Hard to predict	Easy to predict
Bad	Good

So we want **maximum variance reduction**.



WHY USED IN Decision Tree Regression?

Because **Decision Tree Regression** predicts numbers
And prediction works best when numbers inside a leaf are **similar**

Example:

Leaf values = [50, 51, 52]

Prediction = mean = 51 → very accurate

But if leaf values = [10, 80, 25]

Mean = 38.3 → totally wrong

So tree tries splits that **reduce variance**.



ONE-SENTENCE MEANING

Variance Reduction = How much we improved group purity by splitting.



MEMORIZE LIKE THIS

Variance = Messiness / spread of numbers

Variance Reduction = Messiness removed after splitting

Goal = Maximum variance reduction



MINI EXAMPLE

Values:

[10, 100, 11, 95]

Variance is very high.

Split 1 (bad):

[10, 100], [11, 95] → still mixed → no improvement

VR = 0

Split 2 (good):

[10, 11], [95, 100] → values close → big improvement
VR = positive value

So choose split 2.

Conclusion

Term	Real Meaning
Variance	How spread out the numbers are
Variance Reduction	How much spread decreased after splitting
Goal	Make small groups whose values are close