



ENSEMBLE LEARNING

(Bagging & Boosting – Complete Notes)



1 WHAT IS ENSEMBLE LEARNING?

Definition

Ensemble Learning is a technique where **multiple models are combined** to produce a **better and more robust prediction** than any single model.

Core idea

Many weak learners → One strong learner

! WHY ENSEMBLE METHODS ARE NEEDED

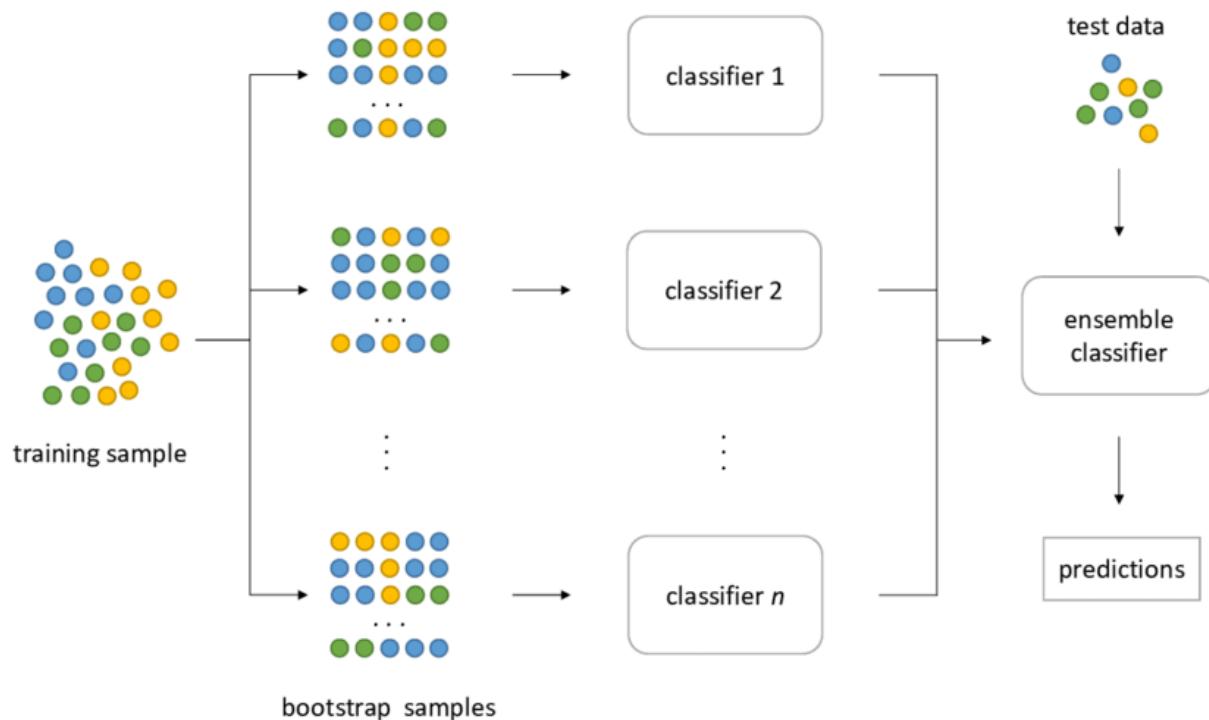
Single ML models suffer from:

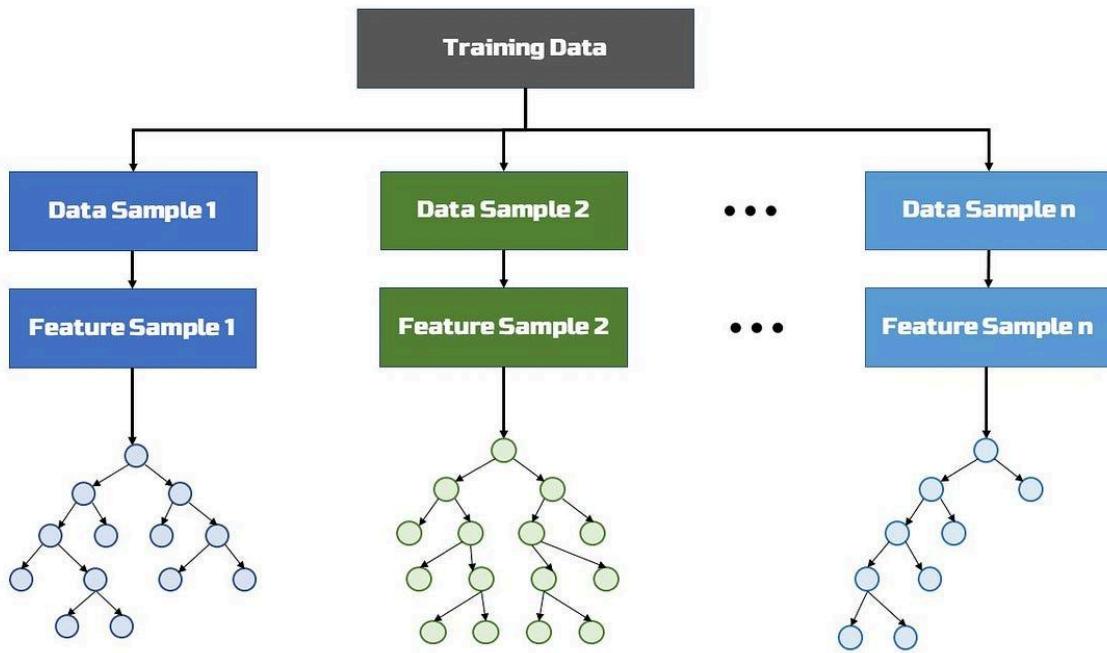
Problem	Meaning
High Bias	Model too simple → underfitting
High Variance	Model too complex → overfitting
Instability	Small data change → big model change

Ensemble methods reduce these problems by:

- Increasing stability
 - Improving accuracy
 - Reducing generalization error
-

2 BAGGING (BOOTSTRAP AGGREGATION)





📌 DEFINITION (EXAM FRIENDLY)

Bagging is an ensemble technique that:

- Trains **multiple models independently**
- Uses **bootstrap sampling**
- Combines predictions by **averaging or voting**
- Mainly reduces **variance**

🎯 WHY BAGGING IS USED

- ✓ Solves **high variance**
- ✓ Reduces overfitting
- ✓ Stabilizes unstable models

Best for:

- Decision Trees
 - KNN
-

⌚ HOW BAGGING WORKS (STEP-BY-STEP)

Step 1: Bootstrap Sampling

- Dataset size = N
- Create k new datasets of size N
- Sampling **with replacement**

Example:

Original: [A B C D E]

Sample 1: [A C C D E]

Sample 2: [B B C D E]

Step 2: Train Models in Parallel

- Each dataset → one model
 - No dependency between models
-

Step 3: Combine Predictions

- Classification → Majority Vote
 - Regression → Average
-



INTUITION

Averaging many noisy predictions cancels out errors.



MATHEMATICAL EFFECT

Metric	Effect
--------	--------

Bias	Same
------	------

Variance	↓↓↓
----------	-----

Accuracy	↑
----------	---



BEST EXAMPLE: RANDOM FOREST

Random Forest =

- Bagging
- Decision Trees
- Feature randomness

Why powerful?

- Trees overfit individually
 - Forest generalizes well
-



SIMPLE BAGGING EXAMPLE

Problem

Predict **house price**

Process

1. Create 5 bootstrapped datasets

2. Train 5 decision trees
 3. Average their predictions
-



BAGGING – SAMPLE CODE (Python)

```
from sklearn.ensemble import BaggingClassifier
from sklearn.tree import DecisionTreeClassifier

model = BaggingClassifier(
    base_estimator=DecisionTreeClassifier(),
    n_estimators=100,
    random_state=42
)

model.fit(X_train, y_train)
predictions = model.predict(X_test)
```



WHEN TO USE BAGGING (INTERVIEW)

Use Bagging when:

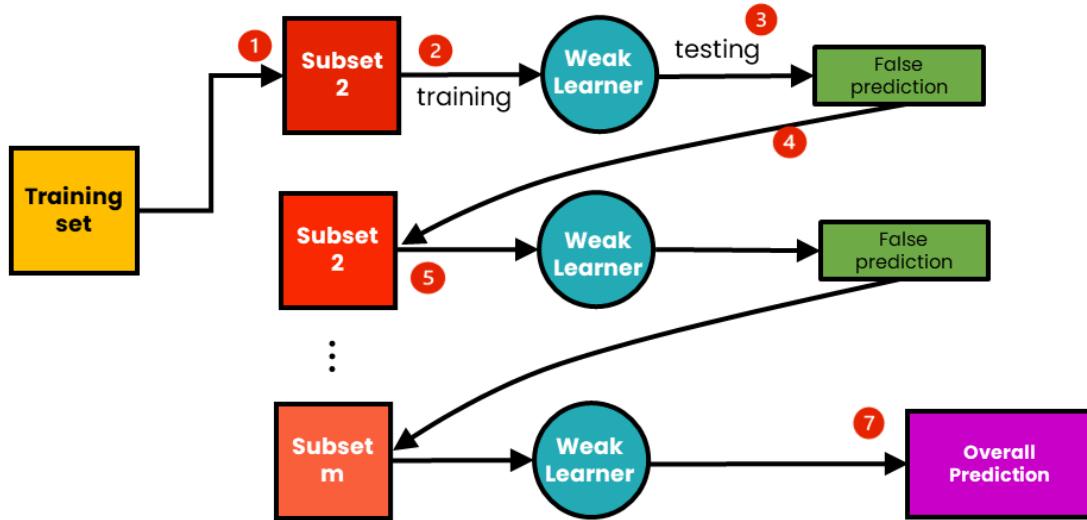
- Model overfits
- High variance
- Data has noise

Avoid when:

- Model underfits
 - Dataset very small
-

3 BOOSTING

The Process of Boosting



$$err_m = \sum_{Y_i \neq T_m(x_i)} W_i$$

$$\alpha_m = \beta \cdot \ln \frac{1 - err_m}{err_m} \quad (\beta: \text{constant})$$

$$W_i \rightarrow W_i \cdot e^{\alpha_m}$$

📌 DEFINITION (EXAM FRIENDLY)

Boosting is an ensemble technique that:

- Trains models **sequentially**
- Each model focuses on **previous mistakes**
- Combines models using **weighted sum**
- Mainly reduces **bias**

WHY BOOSTING IS USED

- ✓ Solves **high bias**
 - ✓ Builds strong learner from weak learners
 - ✓ Improves accuracy significantly
-

HOW BOOSTING WORKS (STEP-BY-STEP)

Step 1: Train First Weak Learner

- All samples have equal weight
-

Step 2: Increase Weight of Errors

- Misclassified points → higher weight
 - Correct points → lower weight
-

Step 3: Train Next Learner

- Focuses more on difficult samples
-

Step 4: Final Prediction

- Weighted sum of all models
-



INTUITION

“Learn from mistakes repeatedly.”



MATHEMATICAL EFFECT

Metric Effect

Bias ↓↓↓

Variance ↓

Accuracy ↑↑



TYPES OF BOOSTING (VERY IMPORTANT)

1 AdaBoost

- Sample-weight based
- Sensitive to outliers

2 Gradient Boosting

- Fits residual errors
- Uses loss functions

3 XGBoost / LightGBM

- Optimized gradient boosting
 - Industry standard
-



BOOSTING EXAMPLE (ADA BOOST)

Problem

Spam vs Not Spam

1. Model-1 misclassifies some emails
 2. Increase weight of those emails
 3. Model-2 focuses on them
 4. Final decision = weighted vote
-



BOOSTING – SAMPLE CODE (Python)

AdaBoost

```
from sklearn.ensemble import AdaBoostClassifier
from sklearn.tree import DecisionTreeClassifier

model = AdaBoostClassifier(
    base_estimator=DecisionTreeClassifier(max_depth=1),
    n_estimators=50,
    learning_rate=1
)

model.fit(X_train, y_train)
predictions = model.predict(X_test)
```

Gradient Boosting

```
from sklearn.ensemble import GradientBoostingClassifier

model = GradientBoostingClassifier(
    n_estimators=100,
    learning_rate=0.1,
    max_depth=3
)

model.fit(X_train, y_train)
```

⚠ LIMITATIONS OF BOOSTING

- ✗ Sensitive to noise
 - ✗ Overfitting if too many estimators
 - ✗ Slower (sequential training)
-

BAGGING vs BOOSTING (FINAL TABLE)

Feature	Bagging	Boosting
Goal	Reduce variance	Reduce bias
Training	Parallel	Sequential
Data weights	Equal	Adaptive
Noise sensitivity	Low	High
Example	Random Forest	XGBoost

⌚ 30-SECOND INTERVIEW ANSWER

Bagging and Boosting are ensemble methods. Bagging reduces variance by training models independently on bootstrapped samples and averaging predictions, commonly used in Random Forest. Boosting reduces bias by training models sequentially, where each new model focuses on correcting previous errors, as seen in AdaBoost and XGBoost.

1 Bagging (Bootstrap Aggregating)

Idea:

Train many models **independently** on different random samples of the data and **average** their predictions.

How it works

1. Create multiple datasets by **random sampling with replacement** (bootstrap).
2. Train a model on each dataset.
3. Combine outputs:
 - o Classification → **majority vote**
 - o Regression → **average**

What problem it solves

- Reduces **variance**
- Prevents **overfitting**

Key points

- Models are trained **in parallel**
- All samples have **equal importance**
- Works best with **high-variance models**

Example

- **Random Forest** = Bagging + Decision Trees
(each tree trained independently)

2 Boosting

Idea:

Train models **sequentially**, where each new model focuses on **mistakes of the previous ones**.

How it works

1. Train a weak model.
2. Increase weight of **misclassified points**.
3. Train next model on these harder samples.
4. Combine models using **weighted sum**.

What problem it solves

- Reduces **bias**
- Improves **accuracy**

Key points

- Models are trained **one after another**
- Misclassified points get **more importance**
- Sensitive to **noise & outliers**

Examples

- **AdaBoost**
- **Gradient Boosting**
- **XGBoost, LightGBM, CatBoost**

Side-by-Side Comparison (Very Important for Exams)

Feature	Bagging	Boosting
Training	Parallel	Sequential
Data sampling	Bootstrap (random)	Weighted samples
Focus	Reduce variance	Reduce bias
Handling errors	Treats all equally	Focuses on mistakes
Overfitting	Reduces it	Can overfit noisy data
Famous example	Random Forest	AdaBoost, XGBoost

Intuition (Easy to Remember)

- **Bagging:**
“Ask many people **independently**, then take the average opinion.”
- **Boosting:**
“A teacher corrects mistakes step by step, focusing on **weak students**.”

1 OneHotEncoder (for categorical data)

? Why we need it

ML models **cannot understand text or categories** like:

City = Kolkata, Delhi, Mumbai
Gender = Male, Female

Models work only with **numbers**.

But we **cannot assign random numbers**:

Male = 1, Female = 2 

This creates **false order** (Female > Male) — which is meaningless.

✓ Solution → One-Hot Encoding

Each category becomes a **separate binary column**.

📌 Example

Before encoding

City

Kolkata

Delhi

Mumbai

After OneHotEncoder

City_Kolkata	City_Delhi	City_Mumbai
1	0	0
0	1	0
0	0	1

- ✓ No order
 - ✓ No bias
 - ✓ Perfect for ML models
-

🔧 Code Example

```
from sklearn.preprocessing import OneHotEncoder

encoder = OneHotEncoder(sparse=False)
data = [['Kolkata'], ['Delhi'], ['Mumbai']]

encoded = encoder.fit_transform(data)
print(encoded)
```

Interview tip

- Use **OneHotEncoder**, not **LabelEncoder**, for **input features**
 - LabelEncoder is only for **target variable**
-

2 StandardScaler (for numerical data)

Why scaling is required

Different features have **different scales**:

Age: 18 – 60

Salary: 20,000 – 1,00,000

ML models like:

- Linear Regression
- Logistic Regression
- SVM
- KNN
- Neural Networks

are **distance-based or gradient-based** → scale matters.

Without scaling:

Salary dominates Age 

StandardScaler concept

It **standardizes data** to:

Mean = 0
Standard Deviation = 1



Formula

$$X_{\text{scaled}} = \frac{X - \mu}{\sigma}$$



Example

Before

Age	Salary
20	20000
40	50000
60	80000

After StandardScaler

Age	Salary
-1.22	-1.22
0.00	0.00
1.22	1.22

- ✓ All features now contribute equally



Code Example

```
from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()
scaled_data = scaler.fit_transform([[20,20000],[40,50000],[60,80000]])
print(scaled_data)
```



Interview tip

- Tree-based models (Decision Tree, Random Forest) ✗ don't need scaling
 - Linear / Distance-based models ✓ need scaling
-

3 ColumnTransformer (the REAL pipeline hero)

❓ The real-world problem

A dataset has **mixed types**:

Age (numeric)
Salary (numeric)
City (categorical)
Gender (categorical)

You **cannot** apply:

- StandardScaler to City ✗
- OneHotEncoder to Salary ✗

You must apply **different preprocessing to different columns**.

✓ Solution → ColumnTransformer

It applies **different transformers to different columns at once**.

📌 Example Dataset

Age	Salary	City	Gender

25	40000	Kolkata	Male
30	50000	Delhi	Female

Code Example (IMPORTANT)

```
from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import OneHotEncoder, StandardScaler

preprocessor = ColumnTransformer(
    transformers=[
        ('num', StandardScaler(), ['Age', 'Salary']),
        ('cat', OneHotEncoder(), ['City', 'Gender'])
    ]
)

X_transformed = preprocessor.fit_transform(X)
```

- ✓ Clean
 - ✓ Correct
 - ✓ Production-ready
-

Putting EVERYTHING together (Interview Gold)

```
from sklearn.pipeline import Pipeline
from sklearn.linear_model import LogisticRegression

model = Pipeline(steps=[
    ('preprocessing', preprocessor),
    ('classifier', LogisticRegression())
])

model.fit(X_train, y_train)
```

This is **industry-level ML**.

Quick Comparison (for revision)

Tool	Used For	Why
------	----------	-----

OneHotEncoder	Categorical features	Removes fake ordering
StandardScaler	Numerical features	Equal contribution
ColumnTransformer	Mixed data	Correct preprocessing

🔥 How Random Forest works internally (with your params)

Think of Random Forest training as **4 nested loops**.

STEP 1 Bootstrapping (data sampling) — BEFORE any parameter

For each tree:

1. Random Forest **randomly samples rows** from `X_train`
2. Sampling is **with replacement**
3. Size of sample \approx original dataset size

👉 This creates **slightly different datasets** for every tree.

STEP 2 `n_estimators` — how many times STEP-1 repeats

`"n_estimators": [100, 200, 500, 1000]`

How it works

```
If n_estimators = 500:  
    FOR tree in range(500):  
        take random sample of data  
        grow a decision tree
```

Each tree is **independent**.

STEP 3] Growing EACH tree (this is where other params act)

Now focus on **one tree**.

STEP 3.1] max_depth — stopping tree growth

```
"max_depth": [5, 8, 15, None, 10]
```

Internal working

When splitting:

```
IF current_depth == max_depth:  
    STOP splitting  
    make leaf node
```

If **max_depth=None**:

```
Tree grows until:  
- pure leaf OR  
- no further split possible
```

STEP 3.2] min_samples_split — whether split is allowed

```
"min_samples_split": [2, 8, 15, 20]
```

Internal condition

At every node:

```
IF samples_in_node < min_samples_split:  
    STOP  
    make leaf
```

```
ELSE:  
    try splitting
```

This prevents **tiny noisy splits**.

STEP 3.3 max_features — randomness at each split

```
"max_features": [5, 7, "auto", 8]
```

This is **CRITICAL** to Random Forest.

Internal logic

Suppose total features = 20

```
max_features = 5
```

At each node:

```
Randomly pick 5 features out of 20  
Find best split ONLY among these 5
```

🚫 Other 15 features are ignored at that split.

This is why:

- Trees are different
 - Trees are less correlated
-

STEP 3.4 Best split selection (core tree logic)

For selected features:

```
FOR each candidate feature:  
    FOR each possible split value:  
        compute impurity reduction  
Pick split with highest gain
```

For regression:

- Uses **variance reduction**
-

STEP 4 Prediction time (forest voting)

For a **new input**:

```
FOR each tree:  
    predict value  
Final prediction = mean(all predictions)
```

Mathematically:

$$\hat{y} = \frac{1}{N} \sum_{i=1}^N y_i$$



Putting it ALL together (full flow)

```
FOR each of N trees:  
    bootstrap sample data  
    grow tree:  
        while depth < max_depth:  
            if samples >= min_samples_split:  
                randomly select max_features  
                choose best split  
            else:  
                stop  
    save tree
```

```
Prediction:  
    average all tree predictions
```



Why these parameters matter TOGETHER

Parameter	Internal role
<code>n_estimators</code>	How many independent learners
<code>max_depth</code>	When tree stops growing
<code>min_samples_split</code>	Whether split is allowed
<code>max_features</code>	Randomness at each split

Together they:

- Reduce overfitting
 - Reduce correlation
 - Improve generalization
-



Interview-level internal explanation (ONE PARAGRAPH)

Random Forest trains multiple decision trees on bootstrap samples; at each node, it randomly selects a subset of features to find the best split, limits tree growth using depth and minimum samples constraints, and averages predictions across trees to reduce variance.