

Got it 

You want **YOUR notes**, cleaned, structured, and **complete**, using **only the information you provided**, with **Step 6 made crystal-clear** and the **final comparison table added**.

No new theory, no confusion.

Below is your **final, exam + revision-ready AdaBoost notes**.

AdaBoost (Adaptive Boosting) — COMPLETE & CLEAN NOTES

1 What is AdaBoost? (Definition)

AdaBoost is a boosting algorithm that combines many weak learners into one strong learner by repeatedly focusing more on the mistakes made earlier.

- **Ada** = Adaptive
- **Boost** = Improve performance by combination

 **Core idea:**

 Hard examples should get more attention.

2 Why AdaBoost is needed (Motivation)

Problem with a single model:

- One decision tree can:
 - Overfit (high variance) 
 - Underfit (high bias) 

AdaBoost's solution:

- Use **very simple models (decision stumps)**
- Combine them **sequentially**
- Each new model focuses on **previous errors**



Goal:

Turn many weak learners into one strong learner.

3 What is a Weak Learner in AdaBoost?

A **weak learner** is a model that performs only slightly better than random guessing.

In AdaBoost:

- Weak learner = **Decision Tree Stump**
- Depth = 1
- Single split (e.g., **Salary ≤ 50K**)



Why stumps?

- High bias
 - Low variance
 - Easy to control
-

4 Big Picture: How AdaBoost Works

Data → DT Stump 1 → DT Stump 2 → DT Stump 3 → ... → Final Model

- Models are trained **one after another**
 - **NOT independently** (unlike Random Forest)
-

5 Step-by-Step Working of AdaBoost

◆ Step 1: Initialize sample weights

If dataset has **N** samples:

Weight of each sample = $1 / N$



Meaning:

Initially, **all data points are equally important**.

◆ Step 2: Train the first weak learner

- Train a **decision stump**
 - Choose the split with **lowest weighted classification error**
 - Use **entropy / gini**
-

◆ Step 3: Calculate Total Error (TE)

TE = sum of weights of misclassified samples

📌 Smaller error → better stump.

◆ Step 4: Compute performance of the stump (α)

$$\alpha = 1/2 \times \ln((1 - TE) / TE)$$

📌 Interpretation:

- Small error → **large α** (important learner)
- Large error → **small α** (less important)

🎯 α decides how much say this stump has in final prediction.

◆ Step 5: Update sample weights

- **Correctly classified samples** → weight decreases
- **Misclassified samples** → weight increases

Mathematically:

Correct: $w \times e^{-\alpha}$

Wrong: $w \times e^{+\alpha}$

📌 Meaning:

👉 “Pay more attention to mistakes next time.”

🔥 Step 6: Normalize the Weights (VERY IMPORTANT)

✓ What normalization means (one line)

Normalization rescales weights so that their total becomes 1 again.

❓ Why Step 6 is required

After Step 5:

- Some weights increase
- Some weights decrease

Example weights:

0.058, 0.058, 0.058, 0.058, 0.349, 0.058, 0.058

Sum:

0.697 ✗ (not equal to 1)

This is a problem.

✗ What goes wrong without normalization

AdaBoost treats weights as **probabilities** when:

- Sampling data
- Selecting samples for the next stump

📌 Probabilities must sum to 1

Without normalization:

- Sampling becomes incorrect
 - Algorithm breaks logically
-

What normalization actually does

`new_weight = weight / (sum of all weights)`

Tiny numeric example

Before normalization:

Weights = [0.058, 0.058, 0.349]

Sum = 0.465

After normalization:

$$0.058 / 0.465 = 0.125$$

$$0.058 / 0.465 = 0.125$$

$$0.349 / 0.465 = 0.75$$

Now:

$$0.125 + 0.125 + 0.75 = 1 \quad \checkmark$$

What this achieves

- Hard sample (0.75) → very likely to be chosen again
- Easy samples (0.125) → less likely to be chosen

 This forces the next stump to focus on mistakes.

Relation to your PDF

- Bins like 0–0.08, 0.08–0.16, ...
- These bins are created after normalization
- Used for random sampling

 Without normalization → bins make no sense.

One-line memory lock

Normalization does not change importance — it only rescales weights so they behave like probabilities.

◆ Step 7: Train next stump

- New stump sees **harder samples more often**
 - Repeat **Steps 2 → 6**
-

◆ Step 8: Final Prediction (**VERY IMPORTANT**)

Final model:

$$f(x) = \alpha_1 h_1(x) + \alpha_2 h_2(x) + \dots + \alpha_m h_m(x)$$

- Each stump **votes**
- Vote strength = α

 **Class with higher total weighted vote wins.**

6 What AdaBoost USES

- ✓ Sample weights
 - ✓ Decision tree stumps
 - ✓ Weighted error
 - ✓ Exponential loss
 - ✓ Weighted voting
-

7 What AdaBoost DOES NOT USE

- ✗ Log-odds
- ✗ Sigmoid
- ✗ Gradients
- ✗ Residuals

-  Explicit learning rate
 -  Second-order derivatives
 -  Regularization parameters
-

ONE TABLE (REVISION GOLD)

Algorithm	Uses log-odds?	Uses sigmoid?	Uses gradients?	Output
AdaBoost	 No	 No	 No	Class (+1 / -1)
Gradient Boosting	 Sometimes	 Sometimes	 Yes	Value / Probability
XGBoost	 Yes	 Yes	 Yes	Probability

8 When should you use AdaBoost?

-  Clean datasets
 -  Small datasets
 -  Conceptual learning
 -  When interpretability matters

 -  Very noisy data
 -  Large-scale production systems
-

One-Line Memory Rules (REVISION GOLD)

- AdaBoost focuses on **data weights**, not gradients
 - Hard samples get **higher weight**
 - Better stump → **higher α**
 - Final decision = **weighted vote**
-

Interview-Ready Definition

“AdaBoost is an ensemble boosting algorithm that sequentially trains weak learners and adapts by increasing the weight of misclassified samples so that subsequent learners focus more on difficult cases.”

Final Mental Model (LOCK THIS)

AdaBoost = Learn → Find mistakes → Focus harder → Combine votes

If you want next:

- Convert this into **1-page exam sheet**
- **Numerical AdaBoost example**
- **AdaBoost vs XGBoost interview Q&A**

Just say 