



AdaBoost (Adaptive Boosting) — Complete Concept-Wise Notes

1 What is AdaBoost?

AdaBoost is an **ensemble learning algorithm** that combines multiple **weak learners** to build a **strong classifier**.

- Mainly used for **binary classification**
 - Commonly uses **decision stumps** as weak learners
 - Works on the principle of **boosting** (learning from mistakes)
-

2 Why AdaBoost is Needed

Single classifiers:

- May have high bias or fail on complex data
- Treat all training samples equally

AdaBoost improves performance by:

- Giving more importance to **hard-to-classify samples**
 - Giving more importance to **accurate learners**
-

3 Key Terminology (Must Know)

Weak Learner

- A simple model slightly better than random guessing
- Usually a **decision stump**

Sample Weight

- Importance assigned to each training data point
- Updated after every iteration

Model Weight (α)

- Importance assigned to each weak learner
 - Depends on the learner's error
 - Fixed once calculated
-

4 Output Representation

For binary classification, AdaBoost uses numeric labels:

- YES → +1
- NO → -1

This allows weighted mathematical combination.

5 Initialization Step

- All training samples are assigned **equal weights**
- No learner is favored initially
- Sum of all sample weights is normalized to 1

6 Sequential Training Process

AdaBoost trains weak learners **one after another**.

Each new learner is trained using the **updated sample weights** from the previous step.

This is why AdaBoost is called **sequential in training**.

7 Training a Weak Learner

Each weak learner:

- Is trained to minimize **weighted classification error**
 - Focuses more on samples with higher weights
 - Produces predictions in the form +1 or -1
-

8 Weighted Error Concept

- Error is computed as the **sum of weights of misclassified samples**
- This error determines how reliable the learner is

Lower error → better learner

Higher error → weaker learner

9 Model Weight (α) — Concept

- α measures the **reliability** of a weak learner
- It is **inversely proportional** to training error

- Lower error results in higher α
 - Higher error results in lower α
 - Once calculated, α remains **fixed**
 - α is used **only during final prediction**
-

10 Sample Weight Updating (Core Boosting Step)

After training each learner:

Misclassified samples

- Their weights are **increased**
- So they become more important

Correctly classified samples

- Their weights are **decreased**
- So they become less important

Purpose:

- Force the next learner to focus on mistakes
-

11 Normalization Step

- After updating, sample weights are normalized
- Ensures the total weight remains constant
- Prevents uncontrolled weight growth

12 Iterative Learning

Steps 7–11 are repeated:

- For a fixed number of learners, or
- Until sufficient performance is achieved

Each iteration improves focus on harder samples.

13 Prediction Phase (Very Important)

- All trained weak learners participate
- Each learner gives:
 - A prediction (+1 or -1)
 - Its fixed weight (α)
- No learner depends on another

AdaBoost is **parallel in prediction**.

14 Final Prediction Rule

Final score is computed as:

$$F(x) = \sum_{i=1}^M \alpha_i \cdot h_i(x)$$

Decision:

- If $F(x) > 0$ $F(x) > 0 \rightarrow \text{YES}$
- If $F(x) < 0$ $F(x) < 0 \rightarrow \text{NO}$

15 Equivalent Interpretation

- Sum all α values of learners predicting YES
- Sum all α values of learners predicting NO
- Class with the **higher total weight** is selected

⚠ This is **not majority voting**

16 Important Clarifications (Exam Traps)

- AdaBoost increases **sample weights**, not model weights
 - Models with higher error get **lower α**
 - Samples with higher error get **higher weight**
 - Training is sequential
 - Prediction is not sequential
-

17 Advantages of AdaBoost

- Converts weak learners into a strong classifier
 - Reduces bias
 - Simple and theoretically strong
 - Often high accuracy on clean data
-

18 Limitations of AdaBoost

- Sensitive to noisy data
 - Outliers may dominate learning
 - Performance drops with mislabeled data
-

19 One-Line Exam Statements

- **Definition:**
AdaBoost is an ensemble learning algorithm that sequentially trains weak learners by reweighting samples and combines them using weighted voting.
 - **Key Idea:**
AdaBoost focuses on misclassified samples and reliable learners.
-

20 Ultra-Short Revision Rules

- Weak learner → prediction + α
- Misclassified sample → higher weight
- Better learner → higher α
- Final decision → weighted sum

◆ FINAL OVERVIEW: RandomizedSearchCV + RandomForest (Estimator Logic)

1 What you define first (search space)

You give **allowed values**, not trials:

```
rf_params = {  
    "max_depth": [5, 8, 15, None, 10],  
    "max_features": [5, 7, "auto", 8],  
    "min_samples_split": [2, 8, 15, 20],  
    "n_estimators": [100, 200, 500, 1000]  
}
```

This means:

- These are the **choices** RandomizedSearchCV is allowed to pick from.
-

2 What n_iter = 100 really means

```
RandomizedSearchCV(..., n_iter=100)
```

Means:

Try 100 random hyperparameter combinations.

Important:

- It samples **with replacement**
- Values can repeat

- Combinations may still differ because other parameters change
-

3 What happens in ONE iteration (very important)

For each of the 100 iterations:

1. RandomizedSearchCV randomly selects:
 - one value for `max_depth`
 - one value for `max_features`
 - one value for `min_samples_split`
 - one value for `n_estimators`
 2. A new **RandomForest estimator** is created with those values.
-

4 How `n_estimators` works inside each iteration

Suppose one iteration picks:

```
n_estimators = 500
```

Then:

- That one Random Forest model builds **500 decision trees**
- Each tree:
 - Uses bootstrap sampling
 - Uses feature randomness

- Grows based on max_depth, min_samples_split, etc.

So:

n_estimators = number of trees inside ONE model

5 Cross-validation happens next (cv=3)

That same forest is evaluated using:

- 3-fold cross-validation
- Train on 2 folds
- Validate on 1 fold
- Average the score

This gives **one performance score** for that iteration.

6 Repeat this process 100 times

- 100 different Random Forests are trained
 - Each forest has:
 - Different hyperparameters
 - Possibly different number of trees
 - Each one is evaluated independently
-

7 Final selection

After all iterations:

- RandomizedSearchCV compares all CV scores
- Picks the hyperparameter set with the **best average score**
- Stores it as:

`random.best_params_`

8 What `best_estimator_` really is

It is:

A fully trained Random Forest model using the best hyperparameters, retrained on the full training dataset.

Not partial.

Not reused.

Fully rebuilt.

9 Very important mental model (lock this)

```
RandomizedSearchCV (n_iter = 100)
|
└── Forest #1 (100 trees)
└── Forest #2 (500 trees)
└── Forest #3 (200 trees)
└── Forest #4 (1000 trees)
|
└── Select best forest
```

🔑 FINAL ONE-LINE SUMMARY (BEST)

`RandomizedSearchCV` tries many Random Forest models with different configurations, where `n_iter` controls how many models are tested and `n_estimators` controls how many trees each model contains, and finally selects the best-performing model using cross-validation.

Your code

```
adaboost_param = {  
    "n_estimators": [50, 60, 70, 80, 90],  
    "algorithm": [ 'SAMME', 'SAMME.R' ]  
}
```

1 Why hyperparameter tuning is needed for AdaBoost

AdaBoost builds many weak learners step-by-step.

Its performance depends heavily on:

- How many learners you use
- How predictions are combined

If these are not chosen well:

- Model may **underfit** (too weak)
- Or **overfit** (too complex)

👉 Hyperparameter tuning finds the **best balance**.

2 n_estimators — WHY this is tuned

What it means

"n_estimators": [50, 60, 70, 80, 90]

- Number of **weak learners (decision stumps)**
- More estimators → model learns more patterns
- Fewer estimators → model stays simple

Why tune it

- Too small → poor learning
- Too large → overfitting + slow training

So we try **multiple values** and let CV decide.

📌 This is the **MOST important AdaBoost parameter**

3 algorithm — WHY this is tuned

```
"algorithm": ['SAMME', 'SAMME.R']
```

AdaBoost has **two variants** in sklearn.

- ◆ **SAMME**

- Uses **discrete class predictions**
- Slower
- Older method
- More stable on noisy data

- ◆ **SAMME.R**

- Uses **probability estimates**
- Faster convergence
- Usually **better accuracy**
- Default in sklearn

Why tune it

Some datasets:

- Work better with probability-based boosting
- Some work better with discrete boosting

So we let tuning decide.

4 Why ONLY these two parameters are often tuned

Because:

- AdaBoost is **simple**
- Most performance gain comes from:
 - Number of learners
 - Boosting method

For **basic ML projects**, these two are enough.

5 What MORE hyperparameters you CAN add (important)

✓ Most important missing one: **learning_rate**

"learning_rate": [0.01, 0.05, 0.1, 0.5, 1.0]

What it does

- Controls **how much each weak learner contributes**
- Smaller → slower but safer learning
- Larger → faster but risk of overfitting

📌 **learning_rate + n_estimators work together**

6 If base learner is a Decision Tree (advanced)

You can also tune the **base estimator**:

`base_estimator = DecisionTreeClassifier(max_depth=1)`

Then tune:

```
"base_estimator__max_depth": [1,2,3]
```

But:

- This is **advanced**
 - Not required for beginners
-

7 BEST practical parameter grid (recommended)

```
adaboost_param = {  
    "n_estimators": [50, 100, 200],  
    "learning_rate": [0.01, 0.1, 1.0],  
    "algorithm": ["SAMME", "SAMME.R"]  
}
```

This is:

- Industry-safe
 - Interview-safe
 - Project-safe
-

8 One-line exam / interview answer

Hyperparameter tuning in AdaBoost is done to find the optimal number of weak learners, learning rate, and boosting algorithm that balance bias, variance, and generalization performance.

What does algorithm = 'SAMME' or 'SAMME.R' really mean?

Think of AdaBoost as:

"Many small models give opinions, then we combine them."

The difference is how each small model gives its opinion.

① SAMME — simple YES / NO opinions

How SAMME works (in plain words)

- Each weak learner says only:
 - Class A
 - Class B
- No confidence, no probability
- Just a **hard decision**

Like a person saying:

"I vote YES."

Characteristics

- Uses **discrete predictions**
 - Learning is **slower**
 - Less sensitive to probability noise
 - Works better when probability estimates are unreliable
-

② SAMME.R — YES/NO + confidence

How SAMME.R works (in plain words)

- Each weak learner says:
 - **Class A with 80% confidence**
 - **Class B with 20% confidence**

Like a person saying:

"I strongly believe YES (80%)."

Characteristics

- Uses **probability estimates**
 - Learns **faster**
 - Usually gives **better accuracy**
 - Default choice in sklearn
-

3 Why does this difference matter?

Some datasets:

- Are clean and well-separated → probabilities are reliable
- Are noisy or overlapping → probabilities can be misleading

So:

- **SAMME.R** works better when probabilities make sense
 - **SAMME** works better when probabilities are unstable
-

4 Why tune between them?

Because you **don't know in advance**:

- Whether probability estimates are reliable
- Which method will generalize better

So instead of guessing:

Let cross-validation decide.

5 Very simple analogy (this will click)

SAMME

- Ask people: “*Which side are you on?*”

SAMME.R

- Ask people: “*Which side are you on, and how confident are you?*”

Both are valid — depending on people.

6 One-line exam-ready answer

SAMME uses hard class predictions, while SAMME.R uses class probability estimates, and tuning selects the method that performs best for a given dataset.

7 Ultra-short memory line

- **SAMME** → decision only
- **SAMME.R** → decision + confidence

That's all.