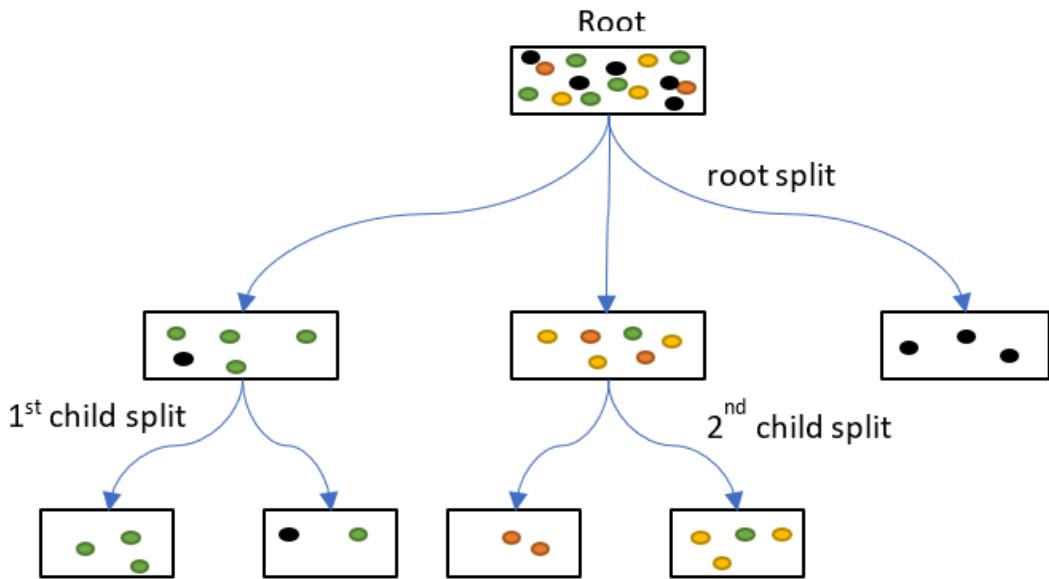


🌳 Decision Tree vs 🌱 Decision Stump



◆ Decision Tree

A **decision tree**:

- Has **multiple splits**
- Can capture **complex patterns**
- High expressive power
- Risk of **overfitting** if deep

Example:

```

Is age > 30?
|--- Yes → Is income > 50k?
|   |--- Yes → Buy
|   |--- No  → Not Buy
|--- No  → Not Buy

```

◆ Decision Stump

A **decision stump**:

- Tree with **depth = 1**
- Only **one split**
- Extremely simple
- Weak learner by design

Example:

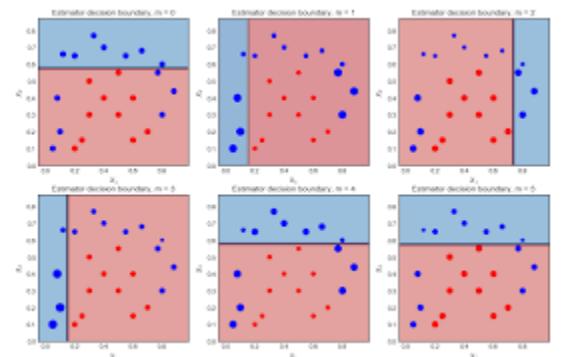
```

Is age > 30?
|--- Yes → Buy
|--- No  → Not Buy

```

🤖 How They Are Used in Boosting

● AdaBoost (Adaptive Boosting)





Default base learner:

Decision Stump

Why stumps?

- AdaBoost works by:
 - Training a weak learner
 - Increasing weight of misclassified points
 - Training another weak learner on hard points
- If learners are **too strong**, boosting becomes unstable

📌 Key idea:

Many **weak rules** + smart weighting = strong classifier

So AdaBoost usually uses:

`depth = 1 (decision stumps)`

➊ Gradient Boosting



Base learner:

Shallow decision trees

Typical depth:

```
depth = 3 to 8
```

Why not stumps only?

- Gradient Boosting fits **residual errors**
- Needs more expressive power
- Captures **feature interactions**

📌 Key idea:

Each tree corrects the mistakes of the previous model



Side-by-Side Comparison

Aspect	Decision Tree	Decision Stump
Depth	>1	1
Complexity	High	Very low
Overfitting	Possible	Almost none
Used in AdaBoost	✗ (rare)	✓ (default)
Used in Gradient Boosting	✓	✗ (usually)

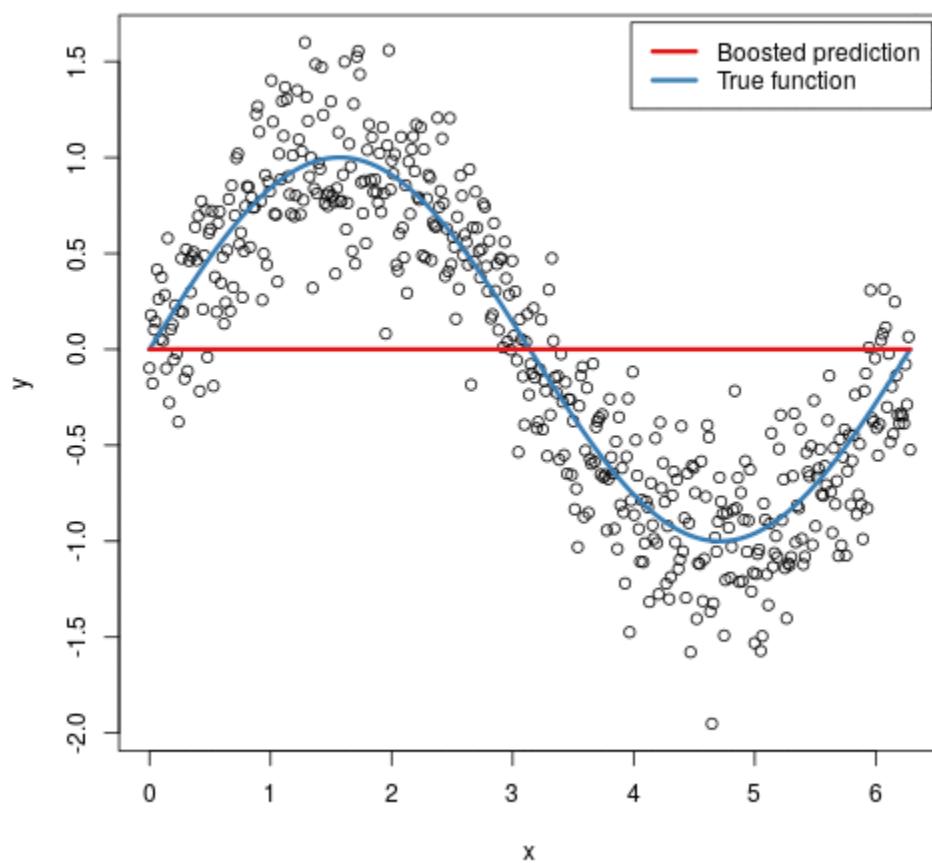
🧠 Why AdaBoost ≠ Gradient Boosting (Important Insight)

AdaBoost	Gradient Boosting
Focuses on misclassified samples	Focuses on residual errors
Sample-weight based	Gradient-based optimization
Works best with stumps	Needs small trees

Sensitive to noise

More robust

GRADIENT BOOSTING — COMPLETE NOTES





1 What is Gradient Boosting? (Core Idea)

Gradient Boosting builds a strong model by adding many weak models sequentially, where each new model learns to correct the errors (residuals) of the previous model using gradient descent.

In short:

- Models are added **one by one**
 - Each model fixes **previous mistakes**
 - Uses **gradients of loss function**
-

2 Why is it called “Gradient” Boosting?

Because it:

- Minimizes a **loss function**
- Uses **gradient descent** logic
- Fits models to the **negative gradient (residuals)**

💡 Residuals = direction in which prediction should move to reduce loss.

3 Why not a single Decision Tree?

Single trees:

- Overfit
- Are unstable

Gradient Boosting:

- Uses **many shallow trees**
 - Each tree is weak
 - Combined model is strong
-

4 What base learner does Gradient Boosting use?

Shallow Decision Trees

Typical:

- Depth = 3 to 8
- NOT decision stumps usually

Why?

- Need some feature interaction
 - Residual patterns are complex
-

5 Mathematical Form (Simple)

The final prediction function is given by:

$$F(x) = F_0(x) + \eta \sum_{m=1}^M h_m(x)$$

Where:

- $F(x)$ is the final strong model
- $F_0(x)$ is the initial model
- $h_m(x)$ is the m-th weak learner
- η is the learning rate (shrinkage parameter)
- M is the total number of weak learners

Where:

- F_0 → initial prediction
 - h_m → m-th tree
 - η → learning rate
 - M → number of trees
-

6 Gradient Boosting Algorithm (Step-by-Step)

STEP 0: Initialize model

For regression:

$F_0 = \text{mean of target values}$

$$F_0 = \text{mean of target values}$$

For classification:

$$F_0 = \log(p / (1 - p))$$

STEP 1: Compute residuals

Residuals are **negative gradient of loss**.

For regression (MSE loss):

$$r_i = y_i - \hat{y}_i$$

STEP 2: Train a tree on residuals

- Fit a **small decision tree**
 - Input = original features
 - Target = residuals
-

STEP 3: Update prediction

- $F_m(x) = F_{\{m-1\}}(x) + \eta * h_m(x)$
 -
 - η = learning rate (0.01 – 0.1)
 -
 -
-

STEP 4: Repeat

Repeat steps 1–3 for **M trees**

7 Worked Example (Regression)

Dataset

$X : [1, 2, 3]$
 $y : [10, 20, 30]$

♦ Initial prediction

$F_0 = \text{mean} = 20$
 $F_0 = \text{mean} = 20$

◆ **Residuals**

$r = [10-20, 20-20, 30-20]$
 $r = [-10, 0, 10]$

◆ **Train Tree₁ on residuals**

Tree predicts:

$h_1(x) = [-8, 1, 7]$

◆ **Update prediction ($\eta = 0.1$)**

New prediction = $20 + 0.1 \times h_1$
= [19.2, 20.1, 20.7]

◆ **Compute new residuals**

$r = [0.8, -0.1, 9.3]$

- Tree₂ fits these residuals
 - Repeat until error is minimized
-

8 Classification Gradient Boosting (Binary)

Uses **Log Loss**:

Log Loss:

$$L = -[y * \log(p) + (1 - y) * \log(1 - p)]$$

Residuals:

$$r_i = y_i - p_i$$

Prediction Update (Sigmoid):

$$p = 1 / (1 + e^{-F(x)})$$

9 Role of Learning Rate (VERY IMPORTANT)

Learning Rate	Effect
Small (0.01)	Slow but stable
Large (0.3)	Fast but risky
Default	0.1



Rule:

Smaller learning rate → more trees needed → better generalization

10 Overfitting Control Parameters

Parameter	Purpose
n_estimators	Number of trees
max_depth	Tree complexity
learning_rate	Step size
subsample	Row sampling
min_samples_leaf	Leaf size
f	

11 Gradient Boosting vs AdaBoost

Aspect	Gradient Boosting	AdaBoost
Error focus	Residuals	Misclassified samples
Math	Gradient descent	Weight update
Noise handling	Better	Sensitive
Base learner	Small trees	Stumps

1 Advantages

- ✓ High accuracy
- ✓ Handles non-linear data
- ✓ Flexible loss functions
- ✓ Works for regression & classification

Gradient Boosting Parameters — Explained Clearly

You wrote:

```
gradient_params = {
    "loss": ['log_loss', 'deviance', 'exponential'],
    "criterion": ['friedman_mse', 'squared_error', 'mse'],
    "min_samples_split": [2, 8, 15, 20],
    "n_estimators": [100, 200, 500],
    "max_depth": [5, 8, 15, None, 10]
}
```

Let's go **one by one**.

1 loss — *WHAT mistake are we trying to reduce?*

Think of **loss** as the judge.

Loss tells Gradient Boosting **what “wrong” means**.

◆ **log_loss / deviance (MOST IMPORTANT)**

Where used?

✓ Classification (binary / multi-class)

What does it measure?

Not just *right vs wrong*, but **how confident you were**.

Example:

Actual	Predicted prob	Loss
Yes	0.51	small
Yes	0.90	very small
Yes	0.10	🔥 huge

👉 Being **confident and wrong** is punished badly.

Why Gradient Boosting loves this

Gradient Boosting:

- Predicts **probabilities**
- Improves predictions gradually

So it needs a loss that says:

“You’re wrong AND overconfident → fix this urgently”

Mental picture

Prediction too confident but wrong



Large loss



Tree focuses strongly on this sample

Why **deviance** = **log_loss**

- Old sklearn name = **deviance**
 - New name = **log_loss**
 - Same meaning, same math
-

◆ **exponential (AdaBoost-style)**

What does it do?

- Heavily increases importance of **misclassified samples**
- Even **one bad outlier** can dominate training

Example:

Correctly classified → small penalty

Wrong → exponentially large penalty

Why risky?

- Noise or label error → model chases it endlessly
- Overfitting risk

📌 That's why:

Gradient Boosting usually prefers **log_loss** over **exponential**

When is **exponential** OK?

- ✓ Clean dataset
 - ✓ Few outliers
 - ✓ Simple classification
-

🎯 SUMMARY for loss

Loss	Meaning	When to use
------	---------	-------------

log_loss	Probability quality	Default, safest
deviance	Same as log_loss	Old name
exponential	Hard punishment	Clean data only

② **criterion** — *HOW does each tree split the data?*

Now imagine the tree is asking:

“Which question should I ask to best reduce error?”

That’s what **criterion** controls.

- ◆ **friedman_mse** (🔥 KEY IDEA)

This criterion is **designed for Gradient Boosting**, not normal trees.

What does it consider?

Not just:

“Does this split make labels pure?”

But:

“Does this split help reduce the overall boosting loss?”

👉 It looks **one step ahead**.

Simple intuition

Normal tree split:

Split to separate Yes / No

Boosting-aware split:

Split to best reduce remaining mistakes

That's the difference.

Why it's best for Gradient Boosting

- Gradient Boosting learns **residuals**
 - Friedman MSE evaluates splits using **residual improvement**
 - Leads to faster convergence & better accuracy
-

◆ **squared_error / mse**

What does it do?

Classic decision tree logic:

- Split to reduce variance
- Ignores boosting context

Why weaker for GB?

- Doesn't consider previous models
- Treats each tree as independent

📌 Works, but not optimal.

🎯 SUMMARY for criterion

Criteri on	Meaning	Use
friedm an_ms e	Boosting -aware split	✓ Best
square d_error	Normal tree split	OK
mse	Same as squared _error	Old name

🧠 Put them together (IMPORTANT)

```
loss = 'log_loss'  
criterion = 'friedman_mse'
```

Means:

"I want to improve probability predictions, and I want each tree split to focus on reducing remaining mistakes."

This is **why Gradient Boosting is powerful**.

3 **min_samples_split** — When is a node allowed to split?

Minimum number of samples required to split a node.

Value	Effect
Small (2)	More splits, complex trees
Large (15–20)	Fewer splits, simpler trees

Controls overfitting:

- Higher value → more regularization
-

4 **n_estimators** — How many trees?

Number of boosting stages (trees).

Trees	Behavior
100	Fast, may underfit
500	Strong, slower
Too many	Overfitting risk

Rule:

More trees = better fit if learning rate is small

5 max_depth — How complex each tree is

Depth of each decision tree.

Depth	Meaning
1	Decision stump
3–5	Best in practice
10+	Risky (overfitting)
None	Full tree (✗ bad for GB)

📌 Important:

Gradient Boosting prefers **shallow trees**, not deep ones.

👉 How ALL these work together (VERY IMPORTANT)

Imagine this setting:

```
n_estimators = 500
max_depth = 5
min_samples_split = 15
loss = log_loss
criterion = friedman_mse
```

Meaning:

- Build **500 small trees**
- Each tree:
 - depth ≤ 5
 - split only if ≥ 15 samples

- Trees focus on **reducing probability error**
- Each tree corrects previous mistakes

→ Result: **Strong but controlled model**

⌚ Common Good Defaults (Interview-safe)

```
GradientBoostingClassifier(  
    loss='log_loss',  
    criterion='friedman_mse',  
    n_estimators=200,  
    max_depth=3,  
    min_samples_split=10,  
    learning_rate=0.1  
)
```

⚠️ Mistakes to Avoid (Important)

- ✗ Using `max_depth=None`
 - ✗ Using deep trees with many estimators
 - ✗ Using `exponential` loss on noisy data
-

🧠 One-line intuition per parameter

Parameter	One-liner
<code>loss</code>	What error am I minimizing?
<code>criterion</code>	How to split nodes inside trees
<code>min_samples_split</code>	When to stop splitting
<code>n_estimators</code>	How many corrections

max_depth

How powerful each correction