



COMPLETE NOTES: K-MEANS, ELBOW METHOD & SILHOUETTE SCORE

(Concept-first • Zero gaps • Exam + Interview ready)

PART 0 — WHY CLUSTERING EXISTS

In Machine Learning, not all datasets have labels.

Examples

- Customers without categories
- Students without grades
- Documents without topics

So we ask:

“Can the data organize itself into meaningful groups?”

This is called **Unsupervised Learning**.

PART 1 — WHAT IS A CLUSTER?

A **cluster** is:

A group of data points that are **more similar to each other** than to points in other groups.

Computers cannot understand “similarity” directly, so we define similarity using **distance**.

PART 2 — DISTANCE (MOST FUNDAMENTAL)

Distance used in K-Means: Euclidean Distance

Formula (copyable):

- $\text{Distance}(x, y) = \sqrt[(x_1 - y_1)^2 + (x_2 - y_2)^2 + \dots]$

Why distance matters:

- Smaller distance → more similar
- Larger distance → less similar

⚠ Distance depends on scale

So:

- Features **must be scaled**
- Otherwise distance becomes meaningless

That's why we use:

- `X_train_scaled`
-

PART 3 — WHAT IS A CENTROID?

A **centroid** is:

The **mean position** of all points in a cluster.

Example

Points:

- $(1, 2), (3, 4), (5, 6)$

Centroid:

- $((1+3+5)/3, (2+4+6)/3) = (3, 4)$

📌 Important:

- Centroid is **not always a real data point**
 - It is a **mathematical average**
-

PART 4 — GOAL OF K-MEANS (KEY STATEMENT)

K-Means places K centroids such that the total squared distance of all points to their nearest centroid is minimized.

That total squared distance has a name.

PART 5 — WCSS / INERTIA (SOUL OF K-MEANS)

Full name:

Within-Cluster Sum of Squares (WCSS)

Meaning:

“How tightly packed are the clusters?”

Mathematical definition (copyable):

- $\text{WCSS} = \sum_{i=1}^K \sum_{x \in C_i} \|x - \mu_i\|^2$

Where:

- C_i = i-th cluster
- μ_i = centroid of cluster i

📌 In scikit-learn:

- `kmeans.inertia_`

So:

- `Inertia` = WCSS
-

PART 6 — WHY SQUARED DISTANCE?

Distance is squared because it:

- Removes negative values
- Penalizes far points more
- Makes optimization stable
- Gives a convex objective per cluster

This is why both **K-Means** and **k-means++** use squared distance.

PART 7 — K-MEANS ALGORITHM (STEP-BY-STEP)

STEP 1: Choose K

- You must specify the number of clusters
- K-Means **cannot decide K itself**

STEP 2: Initialize centroids (k-means++)

Why initialization matters

- Bad initialization → bad clusters

k-means++ logic

1. Choose first centroid randomly
2. Compute distance of each point to nearest centroid
3. Square the distance
4. Choose next centroid with probability \propto squared distance
5. Repeat until K centroids are chosen

📌 Effect:

- Centroids start far apart
- Faster convergence
- Better clustering

STEP 3: Assignment step

For each data point:

- Compute distance to all centroids
- Assign to nearest centroid

STEP 4: Update step

For each cluster:

- Recompute centroid = mean of assigned points
-

STEP 5: Repeat

Repeat steps 3 & 4 until:

- Centroids stop changing
- WCSS stops decreasing

This is called **convergence**.

PART 8 — WHY INERTIA ALWAYS DECREASES

Because:

- More clusters → points closer to centroids
- Squared distances shrink

Extreme cases

- $K = 1 \rightarrow$ very large inertia
- $K = N \rightarrow$ inertia = 0

So inertia is **monotonically decreasing**.

PART 9 — PROBLEM: HOW TO CHOOSE K?

Since inertia always decreases:

- We **cannot** choose K by minimum inertia

Instead, we look for:

Point of diminishing returns

This leads to the **Elbow Method**.

PART 10 — ELBOW METHOD (CONCEPT)

Core idea:

Plot **WCSS vs K** and find the point where improvement slows down.

Before elbow

- Large drop in WCSS
- Meaningful structure captured

After elbow

- Very small improvement
 - Overfitting / noise
-

PART 11 — ELBOW METHOD (CODE)

- `wcss = []`
-
- `for k in range(1, 11):`
- `kmeans = KMeans(n_clusters=k, init="k-means++")`
- `kmeans.fit(X_train_scaled)`
- `wcss.append(kmeans.inertia_)`

Plot:

- `plt.plot(range(1,11), wcss)`
- `plt.xticks(range(1,11))`
- `plt.xlabel("Number of Clusters")`
- `plt.ylabel("WCSS")`
- `plt.show()`

-
- 📌 **wcss** does **not** affect training
 - 📌 It is used **only for analysis**

PART 12 — AUTOMATIC ELBOW DETECTION (KNEELOCATOR)

- `from kneed import KneeLocator`
-
- `k1 = KneeLocator(`
- `range(1,11),`
- `wcss,`
- `curve="convex",`
- `direction="decreasing"`
- `)`
-
- `k1.elbow`

Output:

- `3`

-
- 📌 This automatically finds the elbow point.

PART 13 — SILHOUETTE SCORE (BETTER METRIC)

What it measures

For each point:

- How close it is to its own cluster
- How far it is from other clusters

Silhouette coefficient range:

- **+1** → excellent clustering
 - **0** → overlapping clusters
 - **-1** → wrong clustering
-

Formula (conceptual)

- $s(i) = (b(i) - a(i)) / \max[a(i), b(i)]$

Where:

- $a(i)$ = average distance to own cluster
 - $b(i)$ = average distance to nearest other cluster
-

Code

- `from sklearn.metrics import silhouette_score`
-
- `silhouette_coefficients = []`
-
- `for k in range(2, 11):`
- `kmeans = KMeans(n_clusters=k, init="k-means++")`
- `kmeans.fit(X_train_scaled)`

- `score = silhouette_score(X_train_scaled, kmeans.labels_)`
- `silhouette_coefficients.append(score)`

Example output:

- `[0.5791, ...]`

👉 Starts from `k=2` because silhouette is undefined for `k=1`.

PART 14 — LIMITATIONS (VERY IMPORTANT)

K-Means fails when:

- Clusters are non-spherical
- Cluster sizes differ significantly
- Data contains strong outliers

Elbow method fails when:

- Curve is smooth
 - No clear elbow exists
-

PART 15 — FINAL LOCK-IN SUMMARY

- K-Means minimizes WCSS
- Inertia = WCSS
- Squared distance penalizes far points

- k-means++ improves initialization
- Elbow method selects K using diminishing returns
- KneeLocator automates elbow detection
- Silhouette score measures clustering quality
- Feature scaling is mandatory
-